

Project 2 DESB GRAND CHALLENGE 2015

Juan Gonzalo Quiroz Cadavid¹, Priit Peterson¹, Shivam Maheshwari¹,
and Venkata Narayana Bommanaboina¹

¹University of Tartu

juangonzalo@ut.ee, shivammahe21@gmail.com,
bvnarayana515739@gmail.com, priit.petersonest@gmail.com

March 30, 2025

Overall architecture 1

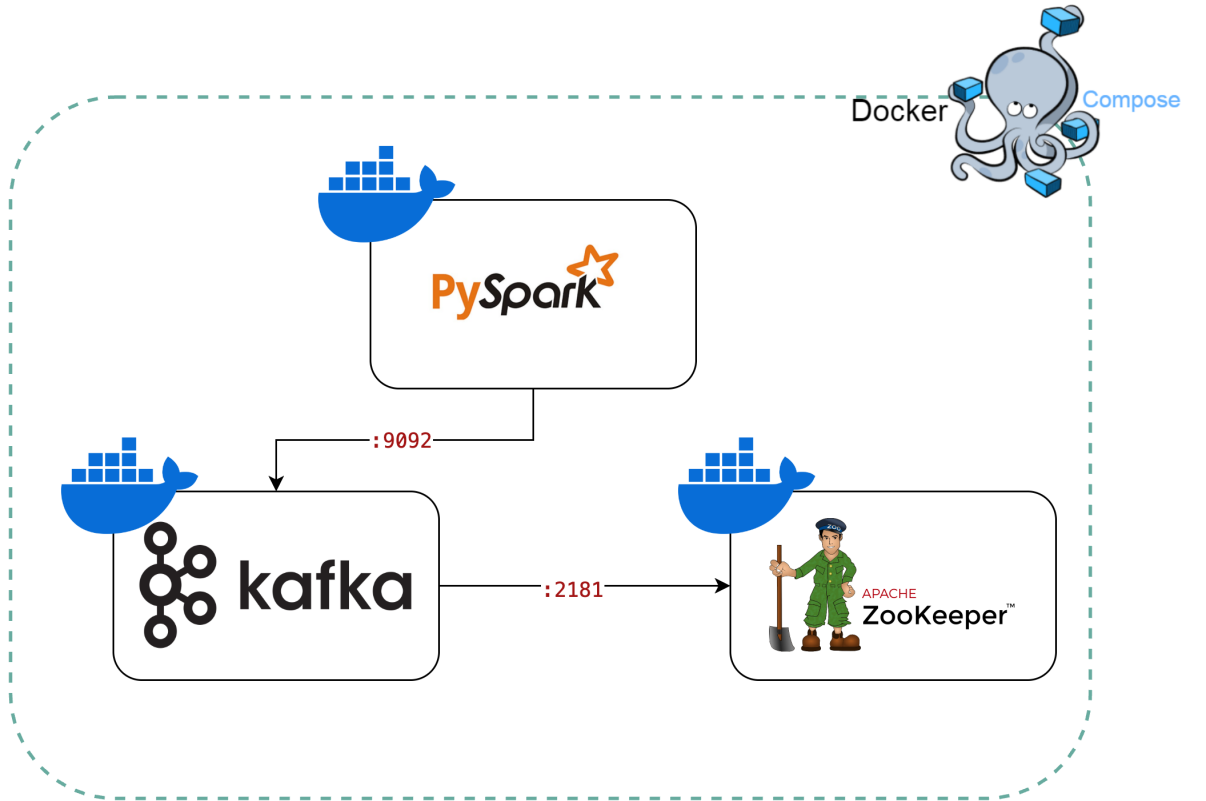


Figure 1: Overall architecture

For this project, the next system architecture were used over docker compose, Kafka, zooKeeper and PySpak as independent containers, communication goes through docker compose networking. Container to container communication was made possible by docker compose networking DNS.

The master plan. 2

- **Data Ingestion & Cleansing**

- Load a portion (e.g., 1 GB) of the NYC Taxi trip dataset.
- Remove invalid rows (e.g., null or zero coordinates, unknown driver IDs) and cast columns to appropriate types.

- **Grid Cell Computation**

- Map each latitude/longitude coordinate to a cell ID.
- For Query 1, use a 500 m grid (cell IDs range from 1.1 to 300.300).
- For Query 2, use a 250 m grid (cell IDs range from 1.1 to 600.600).

- **Query 1: Frequent Routes**

- Compute the 10 most frequent routes in the last 30 minutes, where a route is identified by (`start_cell`, `end_cell`).
- Update results only when the top-10 changes, and output one row containing:
 - * `pickup_datetime`
 - * `dropoff_datetime`
 - * `start_cell_id_n`
 - * `end_cell_id_n`
 - * `delay`

- **Query 2: Profitable Areas**

- Calculate profitability for each area as:

$$\frac{\text{median fare} + \text{tip in last 15 minutes}}{\text{number of empty taxis in last 30 minutes}}$$

- Report the 10 most profitable areas. If fewer than 10 exist, fill in NULL for missing cells. Output one row for each update that includes:
 - * `pickup_datetime`
 - * `dropoff_datetime`
 - * `profitable_cell_id_n`
 - * `empty_taxies_in_cell_id_n`
 - * `median_profit_in_cell_id_n`
 - * `profitability_of_cell_n`
 - * `delay`

- **Streaming Setup & Output**

- Use PySpark Structured Streaming with `foreachBatch` to process micro-batches.
- Maintain global state to detect changes in top-10 routes or profitable areas.
- Write final results to a table or console for each batch, including the delay metric (system time minus ingest time).

2.1 Setup

First, we defined the spark session. We decided to configure spark to work with Delta lake by:

- Adding its JARS package into Spark runtime *[Line 6-7]*.
- Registering Delta SQL commands into SQL extensions *[Line 10-11]*.
- Making it the default Catalog for all tables. *[Line 13-14]*.

Listing 1: Spark session

```
1 builder = SparkSession.builder
2     .appName(
3         "project2_debs_grand_challenge"
4     )
5     .config(
6         "spark.jars.packages",
7         "io.delta:delta-core_2.12:2.4.0"
8     )
9     .config(
10        "spark.sql.extensions",
11        "io.delta.sql.DeltaSparkSessionExtension")
12    .config(
13        "spark.sql.catalog.spark_catalog",
14        "org.apache.spark.sql.delta.catalog.DeltaCatalog"
15    )
16 spark = configure_spark_with_delta_pip(builder).getOrCreate()
```

2.2 Reading and cleaning

We decided to read the CSV without headers, in order to inject the headers as we want to name then on **line 9**, Then, we took 10% of the data, (Around 1Gb) **line 11**

Listing 2: Reading

```
1 columns = [
2     "medallion", "hack_license", "pickup_datetime", "dropoff_datetime",
3     ....
4 ]
5
6 df = spark.read.option("header", "false").csv("data/sorted_data.csv")
7 df = df.toDF(*columns)
8
9 df_sample = df.sample(withReplacement=False, fraction=0.1)
```

After reading, we clean the data by the next operations, which leave us with **13.4 Million elements to analyze**, a before cleaning it was **17.3 Million data**, which means almost **4 Million data was corrupted**

- Changing format of pickup and dropoff date time to timestamp Colum type. 3
- Removing null or 0.0 columns and unknown licenses or drivers. 4
- Convert relevant columns to appropriate data types for numerical computations.

Listing 3: Timestamp

```

1 df_sample = df_sample
2     .withColumn(
3         "pickup_datetime",
4         to_timestamp(
5             col("pickup_datetime"),
6             "yyyy-MM-dd-HH:mm:ss")
7     )
8     .withColumn(
9         "dropoff_datetime",
10        to_timestamp(
11            col("dropoff_datetime"),
12            "yyyy-MM-dd-HH:mm:ss")
13    )

```

Listing 4: Filtering empty fields

```

1
2 df_clean = df_sample.filter(
3     (col("medallion").isNotNull())
4     & (col("medallion") != "0")
5     & (col("medallion") != "UNKNOWN") &
6     (col("hack_license").isNotNull())
7     & (col("hack_license") != "0")
8     & (col("hack_license") != "UNKNOWN") &
9     (col("pickup_datetime").isNotNull()) &
10    (col("dropoff_datetime").isNotNull()) &
11    (col("trip_time_in_secs").isNotNull())
12    & (col("trip_time_in_secs") != 0) &
13    (col("trip_distance").isNotNull())
14    & (col("trip_distance") != 0) &
15    (col("pickup_longitude").isNotNull())
16    & (col("pickup_longitude") != 0.0) &
17    (col("pickup_latitude").isNotNull())
18    & (col("pickup_latitude") != 0.0) &
19    (col("dropoff_longitude").isNotNull())
20    & (col("dropoff_longitude") != 0.0) &
21    (col("dropoff_latitude").isNotNull())
22    & (col("dropoff_latitude") != 0.0) &
23    (col("trip_distance").cast("float") > 0) &
24    (col("fare_amount").cast("float") > 0)
25 )

```

Listing 5: Casting Columns

```

1 df_clean = df_clean
2     .withColumn("trip_time_in_secs", col("trip_time_in_secs").cast("int"))
3     .withColumn("trip_distance", col("trip_distance").cast("float"))
4     .withColumn("fare_amount", col("fare_amount").cast("float"))
5     .withColumn("surcharge", col("surcharge").cast("float"))
6     .withColumn("mta_tax", col("mta_tax").cast("float"))
7     .withColumn("tip_amount", col("tip_amount").cast("float"))
8     .withColumn("tolls_amount", col("tolls_amount").cast("float"))

```

2.3 Query 1

2.3.1 Defining cells and computing

In order to get the cells for each Lat/Lon point, we:

1. Moved the center of the first cell to the corner of the cell (by 0.250 Km to the North, 0.250Km to the West). See Code 6.
2. Shared the new P0.0 across all nodes using Spark context. See Code 6.
3. For upcoming lat/lon attributes, we will only need to calculate the distance to the Point 0 and then divide it by 0.5 Km on both lat and long distances to get the cell location. See Code 7.
4. Defined a new UDF function for calculating the cells locations. See Code 8.
5. Applied the UDF as a new column for all elements in the DF. See Code 9.
6. Filtered by -1_-1, which means the (Lat, Lon) pair is located outside the boundaries. **After cleaning, 54K rows were removed.** See Code 10.
7. Grouped them by tumbling time window of 30 mins, start and end cell and Ordered then by count. See Code 11.
8. **Ta-da!** The resulting top 10 is presented on table 1, Fig 2 shows their distribution as a pie chart.

Listing 6: Moving

```

1 pRoot = [41.474937, -74.913585]
2 p0 = move_point(pRoot[0], pRoot[1], 0.250, 0.250)
3 broadcast_center = spark.sparkContext.broadcast(p0)

```

Listing 7: getCells

```

1 def getCells(center, lat, lon, upTo=300):
2     distLat, distLon = latlon_distance(center[0], center[1], lat, lon)
3     if distLat == -1 or distLon == -1:
4         return -1, -1
5
6     cellLat, cellLon = int(distLat//0.5), int(distLon//0.5)
7
8     if distLat > upTo or distLon > upTo:
9         return -1, -1
10
11    return cellLat, cellLon

```

Listing 8: UDF

```

1 def getCellString(lat, lon):
2     center = broadcast_center.value
3     cellLat, cellLon = getCells(center, lat, lon)
4     return f"{cellLat}_{cellLon}"
5
6 getCellUDF = F.udf(getCellString, StringType())

```

Listing 9: Adding cells

```

1 df_routes = df_clean.withColumn(
2     "start_cell",
3     getCellUDF(col("pickup_latitude"), col("pickup_longitude"))
4 )
5 ).withColumn(
6     "end_cell",
7     getCellUDF(col("dropoff_latitude"), col("dropoff_longitude"))
8 )

```

Listing 10: filterign

```

1 df_routes_cleaned = df_routes.filter(
2     (col("end_cell") != "-1_-1") &
3     (col("start_cell") != "-1_-1")
4 )

```

	Time_window	start_cell	end_cell	Number_of_Rides
0	(('2013-01-15 09:00:00'), ('2013-01-15 09:30:00'))	159_154	159_151	15
1	(('2013-01-14 08:00:00'), ('2013-01-14 08:30:00'))	159_154	160_151	13
2	(('2013-01-15 06:30:00'), ('2013-01-15 07:00:00'))	160_153	159_154	13
3	(('2013-06-20 07:00:00'), ('2013-06-20 07:30:00'))	160_153	160_155	12
4	(('2013-03-07 06:30:00'), ('2013-03-07 07:00:00'))	160_153	159_154	11
5	(('2013-04-15 07:00:00'), ('2013-04-15 07:30:00'))	160_153	160_155	11
6	(('2013-05-07 06:30:00'), ('2013-05-07 07:00:00'))	160_153	159_155	11
7	(('2013-11-14 06:30:00'), ('2013-11-14 07:00:00'))	160_153	160_155	11
8	(('2013-01-14 08:00:00'), ('2013-01-14 08:30:00'))	159_154	159_151	11
9	(('2013-03-13 08:30:00'), ('2013-03-13 09:00:00'))	161_155	164_153	11

Table 1: Top 10 rides

Listing 11: Group by Time and cells

```

1 df_frequent_routes_by_time_window = (
2     df_routes_cleaned
3     .groupBy(
4         F.window("pickup_datetime", "30-minutes"),
5         "start_cell",
6         "end_cell"
7     )
8     .count()
9     .withColumnRenamed("count", "Number_of_Rides")
10    .withColumnRenamed("window", "Time_window")
11 )
12 top10_routes_by_time_window = df_frequent_routes_by_time_window
13    .orderBy(
14        col("Number_of_Rides").desc()
15    )
16    .limit(10)

```




Figure 2: Top 10 rides distribution

2.4 Part 2

To address query 2, we:

- Created a Stream dataframe by writing into disk and loading it as a Stream. See code 12.
- Using the stream, we run a writeStream process on batches. See code 13.
- On every batch, we filter the data for only the last 30 mins, base on their drop off time. See code 14.
- aggregated the output into a single table. See code 15.
- Tables 2, and 3 shows the final result.

Listing 12: Transforming into stream

```

1 df_routes_cleaned
2     .write
3     .format("delta")
4     .mode("overwrite")
5     .save("./output/delta/taxi_data")
6
7 df_stream = spark.readStream
8     .format("delta")
9     .load("./output/delta/taxi_data")

```

Listing 13: Stream

```

1
2 df_stream = df_stream.withColumn(
3     "pickup_datetime",
4     to_timestamp(
5         col("pickup_datetime"),
6         "yyyy-MM-dd-HH:mm:ss"
7     )
8 )
9 df_stream = df_stream.withColumn(
10    "dropoff_datetime",
11    to_timestamp(
12        col("dropoff_datetime"),
13        "yyyy-MM-dd-HH:mm:ss"
14    )
15 )
16 df_stream = df_stream.withColumn("ingest_time", current_timestamp())
17
18 query = (
19     df_stream.writeStream
20     .trigger(once=True)
21     .foreachBatch(process_batch)
22     .outputMode("append")
23     .start()
24 )
25 query.awaitTermination()

```

Listing 14: Bacth operation

```

1
2 def process_batch(batch_df, batch_id):
3     # Skip empty batches
4     if batch_df.rdd.isEmpty():
5         return
6
7     max_dropoff = batch_df.agg({"dropoff_datetime": "max"}).collect()[0][0]
8     if max_dropoff is None:
9         return
10    ref_time = max_dropoff - timedelta(minutes=30)
11
12    # Compute grid cell IDs for pickup and dropoff using the 500m grid
13    ...
14    Calculating cells ids
15    ...
16
17    # Filter for trips with dropoff_datetime >= ref_time (last 30 minutes)
18    df_last30 = batch_df.filter(col("dropoff_datetime") >= F.lit(ref_time))
19    print(f"Window filter: dropoff_datetime >= {ref_time}")
20    print(f"df_last30 count = ", df_last30.count())
21    df_last30.show(5)
22
23    # Aggregate routes and get top 10 most frequent
24    df_frequent_routes = df_last30.groupBy("start_cell", "end_cell") \
25        .count() \
26        .withColumnRenamed("count", "Number_of_Rides")
27    top10_routes = df_frequent_routes
28        .orderBy(
29            col("Number_of_Rides").desc()
30        )
31        .limit(10)
32    top10_list = top10_routes.collect()
33    ...
34    ...
35    for i in range(10):
36        if i < len(top10_list):
37            route = top10_list[i]
38            output_row[f"start_cell_id_{i+1}"] = route["start_cell"]
39            output_row[f"end_cell_id_{i+1}"] = route["end_cell"]
40        else:
41            output_row[f"start_cell_id_{i+1}"] = None
42            output_row[f"end_cell_id_{i+1}"] = None

```

Listing 15: Output

```

1 # This is your foreachBatch function to process each micro-batch
2 def process_batch(batch_df, batch_id):
3     ...
4     ...
5     ...
6     # Define the output schema explicitly
7     output_schema = StructType([
8         StructField("pickup_datetime", TimestampType(), True),
9         StructField("dropoff_datetime", TimestampType(), True),
10        StructField("start_cell_id_1", StringType(), True),
11        ...
12        ...
13        ...
14    ])
15    result_df = spark.createDataFrame([output_row], schema=output_schema)
16    result_df.write.mode("append").saveAsTable("frequent_routes")

```

Pickup Datetime	Dropoff Datetime	Delay (s)
2013-12-31 23:54:08	2014-01-01 00:36:35	58.90

Table 2: Trip Timing Information

Start Cell	End Cell	Start Cell	End Cell
156.160	152.167	154.159	172.160
158.151	166.149	153.163	171.135
161.145	156.162	157.159	158.153
176.157	157.180	153.165	155.158
153.158	154.162	156.161	154.159

Table 3: Top 10 Frequent Routes

Query 2: Profitable Areas 3

In this query, we compute which cells (or areas) offer the highest current profit opportunity. Profitability is defined as

$$\text{Profitability} = \frac{\text{median(fare + tip in last 15 min)}}{\text{number of empty taxis in last 30 min}}.$$

A taxi is considered *empty* in a cell if its last dropoff happened there in the preceding 30 minutes and there has been no subsequent pickup. The system reports the top ten areas at each update, placing them in columns such as `profitable_cell_id_1` through `profitable_cell_id_10`. If fewer than ten cells meet the criteria, the remaining columns are NULL.

3.1 Part 1: 500 m Grid

Listing 16 shows how each micro-batch of streaming data is processed to find the most profitable 500m cells. The code begins by determining the latest `dropoff_datetime` in the batch (`max_dropoff`) and setting up two time windows: one of 15 minutes to calculate the median profit, and another of 30 minutes to count how many taxis remain empty in each cell. In the snippet, `profit_df` uses the `pickup_longitude` and `pickup_latitude` columns to assign each trip to a pickup cell, which is then grouped to find the median of `fare_amount + tip_amount`. Meanwhile, a second aggregation locates each taxi's most recent dropoff cell (within 30 minutes) in order to calculate the number of currently empty taxis. After a join merges these values, the system computes `profitability = median_profit/empty_taxis`, sorts in descending order, then extracts the top ten cells as a single output row. A `delay` is appended to measure the processing latency between reading the triggering event and writing the table row.

Listing 16: Query 2, Part 1: 500 m Grid

```
1 grid_origin_lat = 41.474937
2 grid_origin_lon = -74.913585
3 delta_lat = 0.0045    # about 500m in latitude
4 delta_lon = 0.0060    # about 500m in longitude
5
6 def process_batch_query2_part1(batch_df, batch_id):
7     if batch_df.rdd.isEmpty():
8         print(f"Batch {batch_id}: -empty, -skipping.")
9         return
10
11     max_dropoff = batch_df.agg({"dropoff_datetime": "max"}).collect()[0][0]
12     if not max_dropoff:
13         return
14
15     # Past 15 minutes for profit, past 30 for empty taxis
16     ref_time_profit = max_dropoff - timedelta(minutes=15)
17     ref_time_empty = max_dropoff - timedelta(minutes=30)
18
19     # Build a DF for profit computations
20     profit_df = (
21         batch_df
22         .filter(F.col("dropoff_datetime") >= ref_time_profit)
23         .withColumn("profit", F.col("fare_amount") + F.col("tip_amount"))
```

```

24         .withColumn(
25             "pickup_cell_east",
26             F.floor((F.col("pickup_longitude") - grid_origin_lon) / delta_lon) + 1
27         )
28         .withColumn(
29             "pickup_cell_south",
30             F.floor((grid_origin_lat - F.col("pickup_latitude")) / delta_lat) + 1
31         )
32         .withColumn(
33             "pickup_cell",
34             F.concat(
35                 F.col("pickup_cell_east").cast("int"),
36                 F.lit("."),
37                 F.col("pickup_cell_south").cast("int")
38             )
39         )
40     )
41     profit_agg = profit_df.groupBy("pickup_cell").agg(
42         F.expr("approx_percentile(profit, -0.5) as median_profit")
43     )
44
45     # For each taxi, identify the cell of its most recent dropoff
46     w = Window.partitionBy("medallion").orderBy(F.col("dropoff_datetime").desc())
47     last_dropoff_df = (
48         batch_df
49         .withColumn("rn", F.row_number().over(w))
50         .filter(F.col("rn") == 1)
51         .filter(F.col("dropoff_datetime") >= ref_time_empty)
52         .withColumn(
53             "dropoff_cell_east",
54             F.floor((F.col("dropoff_longitude") - grid_origin_lon) / delta_lon) + 1
55         )
56         .withColumn(
57             "dropoff_cell_south",
58             F.floor((grid_origin_lat - F.col("dropoff_latitude")) / delta_lat) + 1
59         )
60         .withColumn(
61             "dropoff_cell",
62             F.concat(
63                 F.col("dropoff_cell_east").cast("int"),
64                 F.lit("."),
65                 F.col("dropoff_cell_south").cast("int")
66             )
67         )
68     )
69     empty_agg = last_dropoff_df.groupBy("dropoff_cell").agg(
70         F.countDistinct("medallion").alias("empty_taxis")
71     )
72
73     area_df = (
74         profit_agg
75         .join(empty_agg, profit_agg.pickup_cell == empty_agg.dropoff_cell, "inner")

```

```

76         .select(profit_agg.pickup_cell.alias("cell_id"), "median_profit", "empty_
77         .filter(F.col("empty_taxis") > 0)
78         .withColumn("profitability", F.col("median_profit") / F.col("empty_taxis")
79         .orderBy(F.col("profitability").desc())
80     )
81     top10_rows = area_df.limit(10).collect()
82
83     # Determine the event that triggered this output
84     trigger_row = batch_df.orderBy(F.col("dropoff_datetime").desc()).limit(1).collect()
85     pickup_ts    = trigger_row["pickup_datetime"]
86     dropoff_ts   = trigger_row["dropoff_datetime"]
87     ingest_time  = trigger_row.get("ingest_time", dropoff_ts)
88     delay        = (datetime.now() - ingest_time).total_seconds()
89
90     # Build a single output row with up to 10 columns for the profitable cells
91     result_data = {
92         "pickup_datetime": pickup_ts,
93         "dropoff_datetime": dropoff_ts,
94         "delay": delay
95     }
96     for i in range(10):
97         if i < len(top10_rows):
98             info = top10_rows[i]
99             result_data[f"profitable_cell_id_{i+1}"] = info["cell_id"]
100             result_data[f"empty_taxies_in_cell_id_{i+1}"] = info["empty_taxis"]
101             result_data[f"median_profit_in_cell_id_{i+1}"] = info["median_profit"]
102             result_data[f"profitability_of_cell_{i+1}"] = info["profitability"]
103         else:
104             result_data[f"profitable_cell_id_{i+1}"] = None
105             result_data[f"empty_taxies_in_cell_id_{i+1}"] = None
106             result_data[f"median_profit_in_cell_id_{i+1}"] = None
107             result_data[f"profitability_of_cell_{i+1}"] = None
108
109     # (Remaining code: building a schema, writing result_data to a Spark table, etc.)
110
111     df_stream = df_stream.withColumn("ingest_time", F.current_timestamp())
112
113     query2_part1 = (
114         df_stream
115         .writeStream
116         .trigger(once=True)
117         .foreachBatch(process_batch_query2_part1)
118         .outputMode("append")
119         .start()
120     )
121
122     query2_part1.awaitTermination()

```

Whenever a batch includes rides for multiple cells, several columns can be filled. If only one cell meets the criteria, the rest appear as NULL. This behavior is normal because the specification requires that unoccupied columns be blank.

3.2 Part 2: 250 m Grid

Part 2 uses the same reasoning, but cell coordinates are calculated with half the deltas so that each cell measures about 250 m. Listing 17 shows the difference. The same time windows apply: 15 minutes for median profit, 30 minutes for empty taxi counts, and top-ten output columns that fill with NULL if fewer than ten exist. Table 4 illustrates an example row, where `profitable_cell_id_1` indicates the highest profitability and `profitable_cell_id_2` the second highest, etc.

Listing 17: Query 2, Part 2: 250 m Grid

```

1 new_delta_lat = 0.0045 / 2 # 0.00225
2 new_delta_lon = 0.0060 / 2 # 0.0030
3
4 def process_batch_query2_part2(batch_df, batch_id):
5     if batch_df.rdd.isEmpty():
6         print(f"Batch {batch_id}: -empty, -skipping.")
7         return
8
9     max_dropoff = batch_df.agg({"dropoff_datetime": "max"}).collect()[0][0]
10    if not max_dropoff:
11        return
12
13    ref_time_profit = max_dropoff - timedelta(minutes=15)
14    ref_time_empty = max_dropoff - timedelta(minutes=30)
15
16    # The rest is analogous to Part 1, but with new_delta_lat/new_delta_lon
17    ...

```

Column	Example Value
<code>pickup_datetime</code>	2013-12-31 23:54:08
<code>dropoff_datetime</code>	2014-01-01 00:36:35
<code>profitable_cell_id_1</code>	317.319
<code>empty_taxies_in_cell_id_1</code>	1
<code>median_profit_in_cell_id_1</code>	16.0
<code>profitability_of_cell_1</code>	16.0
<code>profitable_cell_id_2</code>	NULL
...	...
<code>delay</code>	42.4519

Table 4: Sample row from the Query 2 (Part 2) output using a 250 m grid.

When only a single cell qualifies, the remaining nine columns are NULL. This is the intended behavior whenever fewer than ten distinct profitable cells are found in the batch and we only used 10% of the whole dataset.